


# DD2417

## 9a: The Transformer architecture

Johan Boye, KTH

# Alignment/Attention mechanism



	la	voiture	verte
the	0.88	0.01	0.01
green	0.01	0.01	0.89
car	0.11	0.98	0.10

# Dependencies in text

Tom visited his brother although he didn't like him much.

# Dependencies in text

Tom visited his brother although he didn't like him much.

Although he was sick, Tom went to work.

# Self-attention

A spring broke in the bicycle.



He went out in the spring sun.




Vaswani et al (2017) presented an alternative translation scheme, the *Transformer* architecture:

Two key ideas:

- Self-attention: Producing contextualized representations of the input words.
- No RNNs: An architecture that allows for parallelization of the training

# Contextualized representations



# Self-attention

	the	spring	sun	shone
the	0.20	0.20	0.55	0.05
spring	0.10	0.40	0.3	0.20
sun	0.10	0.15	0.35	0.40
shone	0.10	0.20	0.40	0.30



# Self-attention

	the	spring	sun	shone	SUM
the	0.20	0.20	0.55	0.05	1
spring	0.10	0.40	0.3	0.20	1
sun	0.10	0.15	0.35	0.40	1
shone	0.10	0.20	0.40	0.30	1

# Self-attention

	the	spring	sun	shone
the	0.20	0.20	0.55	0.05
spring	0.10	0.40	0.30	0.20
sun	0.10	0.15	0.35	0.40
shone	0.10	0.20	0.40	0.30

How much of "sun" do we want to incorporate into the representation of "spring"?

# Self-attention

	$x_1$	$x_2$	$x_3$	$x_4$
$x_1$	$\alpha_{11}$	$\alpha_{12}$	$\alpha_{13}$	$\alpha_{14}$
$x_2$	$\alpha_{21}$	$\alpha_{22}$	$\alpha_{23}$	$\alpha_{24}$
$x_3$	$\alpha_{31}$	$\alpha_{32}$	$\alpha_{33}$	$\alpha_{34}$
$x_4$	$\alpha_{41}$	$\alpha_{42}$	$\alpha_{43}$	$\alpha_{44}$

How much of  $x_3$  do we want to incorporate into the representation of  $x_2$ ?

# Self-attention

$$y_1 = \alpha_{11}x_1 + \alpha_{12}x_2 + \alpha_{13}x_3 + \alpha_{14}x_4$$

$$y_2 = \alpha_{21}x_1 + \alpha_{22}x_2 + \alpha_{23}x_3 + \alpha_{24}x_4$$


$$y_3 = \alpha_{31}x_1 + \alpha_{32}x_2 + \alpha_{33}x_3 + \alpha_{34}x_4$$

$$y_4 = \alpha_{41}x_1 + \alpha_{42}x_2 + \alpha_{43}x_3 + \alpha_{44}x_4$$



	$x_1$	$x_2$	$x_3$	$x_4$
$x_1$	$\alpha_{11}$	$\alpha_{12}$	$\alpha_{13}$	$\alpha_{14}$
$x_2$	$\alpha_{21}$	$\alpha_{22}$	$\alpha_{23}$	$\alpha_{24}$
$x_3$	$\alpha_{31}$	$\alpha_{32}$	$\alpha_{33}$	$\alpha_{34}$
$x_4$	$\alpha_{41}$	$\alpha_{42}$	$\alpha_{43}$	$\alpha_{44}$

# Self-attention



Self-attention  
(first attempt)

# Self-attention: scores

We want to compute a score expressing relevant  $x_i$  is to  $x_j$ .

Idea from w2v and Glove: Score is computed by the dot product:

$$\text{score}(x_i, x_j) = x_i \cdot x_j$$


The *higher* the score, the *more relevant*  $x_j$  is to  $x_i$

- and the more we want to incorporate  $x_j$  into  $y_i$ , the contextualized representation of  $x_i$

**However:** dot product is symmetric, but  $x_j$  might be more relevant to  $x_i$  than the other way around.

# Queries, keys, and values

Words are being used in three roles: As *queries*, *keys*, and *values*.



# Queries, keys, and values

A word vector  $x_i$  is being used in three different roles:

- as an input (a key)
- as the focus of attention when compared to all inputs  $x_1 \dots x_n$  (a query)
- when computing the contextualized output vectors (a value)

Therefore, from each input vector  $x_i$  we will produce three vectors:

- a key vector  $k_i$
- a query vector  $q_i$
- a value vector  $v_i$



# Queries, keys, and values

We define three trainable matrices  $W^Q$ ,  $W^K$ ,  $W^V$  such that:


- $W^Q$ ,  $W^K$  have dimensions  $d \times d_k$ , where  $d$  is the dimensionality of the input vectors  $x_i$
- $W^V$  has dimensions  $d \times d$
- $q_i = W^Q x_i$ ,  $k_i = W^K x_i$ ,  $v_i = W^V x_i$

We can now redefine the score function into:  $\text{score}(x_i, x_j) = \frac{q_i \cdot k_j}{\sqrt{d_k}}$


(we divide by  $\sqrt{d_k}$  to avoid getting too big values)

and compute the proportions using softmax:  $\alpha_{ij} = \frac{\exp(\text{score}(x_i, x_j))}{\sum_k \exp(\text{score}(x_i, x_k))}$


# Queries, keys, and values




# Self-attention



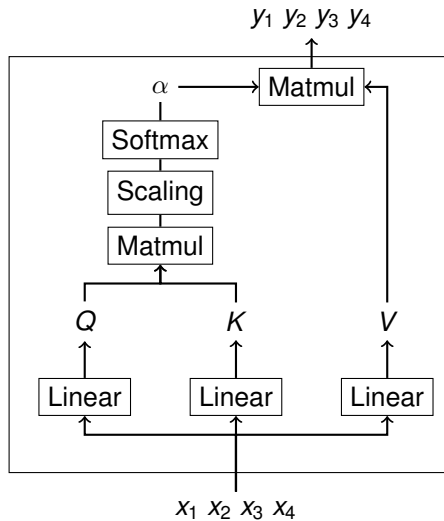
# Self-attention



# Self-attention



# Self-attention



⇐ Self-attention  
(second attempt)

# Self-attention


As an equation, self-attention can be expressed:

$$\text{Self-Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^{\top}}{\sqrt{d_k}}\right)V$$

Vaswani et al. (2017) "Attention is all you need"

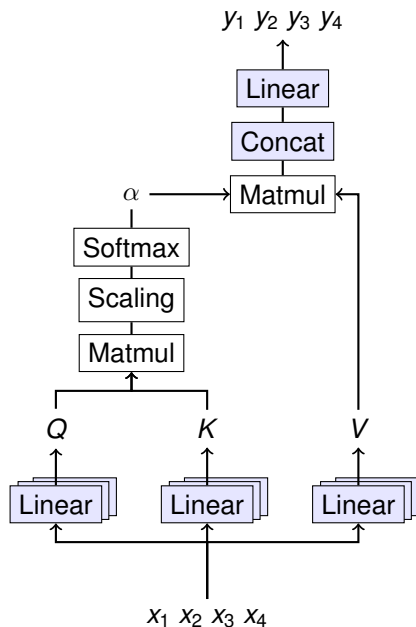
# Multi-head self-attention

A word can need to “pay attention” to many other words, for different reasons.

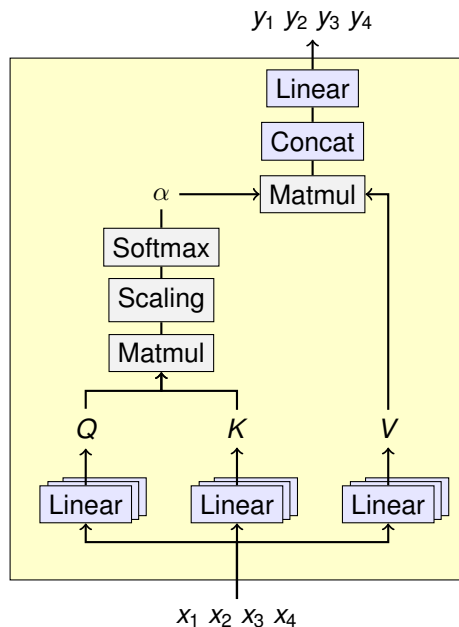




# Multi-head self-attention




# Multi-head self-attention




⇐ Self-attention  
(third attempt)

# Transformer block



# Transformer encoder



# Word order

	the	spring	sun	shone
the	0.20	0.20	0.55	0.05
spring	0.10	0.40	0.3	0.20
sun	0.10	0.15	0.35	0.40
shone	0.10	0.20	0.40	0.30


Word order is not represented in this table, but word order matters!

“Mary is smarter than Tom.”

“Tom is smarter than Mary.”


# Positional encoding

To inject word-order information, each input vector  $x_i$  is added to a *positional vector*  $p_i$  before applying the transformer encoder (or decoder).



Positional vectors could either be *learned* or *generated by a function*.

# Original transformer architecture



# DD2417

## 9b: Masked language models and BERT

Johan Boye, KTH




- BERT – Bidirectional Encoder Representations from Transformers
- A general language model based on the Transformer encoder
- Pre-trained on two objectives:
  - Unmasking words
  - Next sentence prediction
- Through *finetuning* BERT, one can solve a number of language engineering problem.

# Unmasking words

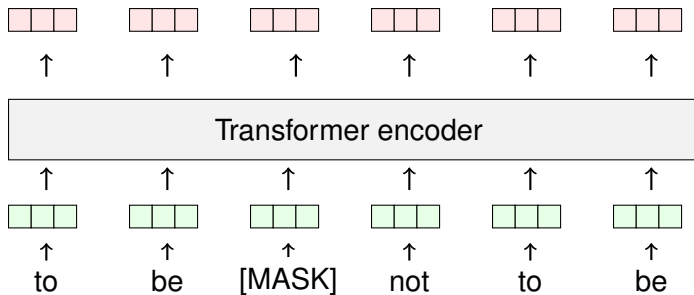
A fill-in-the-gap test:

To be \_\_\_\_\_ not to be

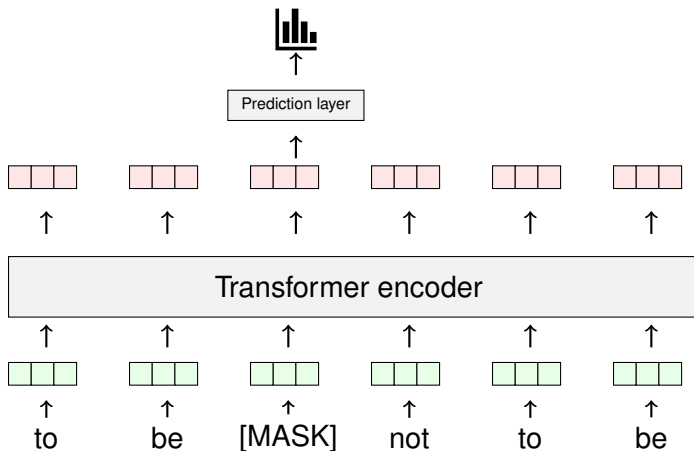
# Unmasking words



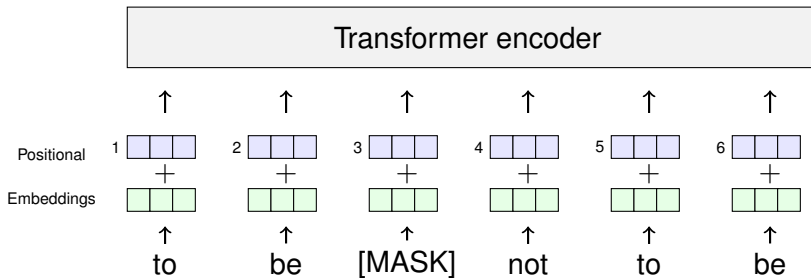
# Unmasking words



# Unmasking words



# Adding positional vectors



# Subword tokenization

Treating every unique word as a unique token class has its drawbacks:

- Some words are rare
  - (e.g. *ultracrepidarian*)
- Similar-looking words can mean similar things
  - (e.g. *contraction*–*contractions*)

Thus, *subword tokenization* makes a lot of sense.

# Subword tokenization

ultracrepedarian  $\rightarrow$  'ultra', '##cre', '##ped', '##arian'

contractions  $\rightarrow$  'contraction', '##s'

debug  $\rightarrow$  'de', '##bu', '##g'

BERT uses 30,000 token classes, computed automatically.

The *byte-pair encoding* (BPE) algorithm can be used (ch 2.4.3 in the textbook).



# Next sentence prediction

Understanding the logical order of sentences:

*Stockholm is the capital of Sweden. One million people live there.*

is OK, but not

*One million people live there. Stockholm is the capital of Sweden.*

# Adding segments


[CLS]

Stockholm ... Sweden

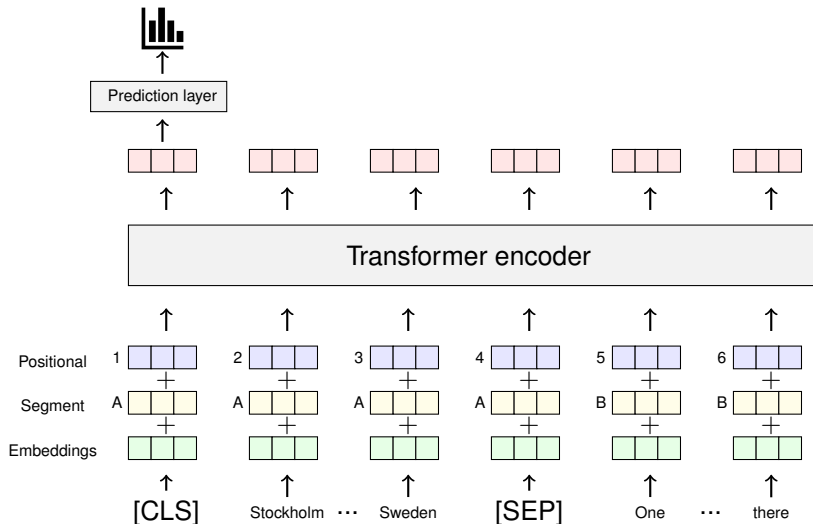
[SEP]

One ... there

# Adding segments



# Adding segments

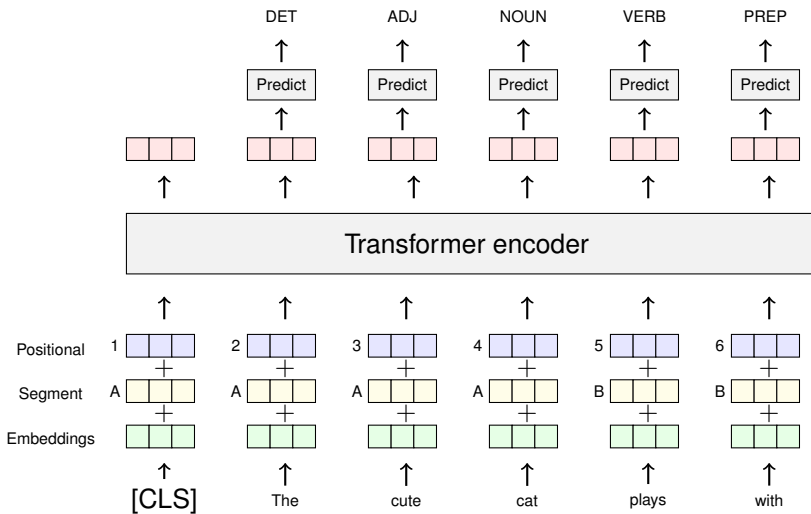


# Fine-tuning BERT

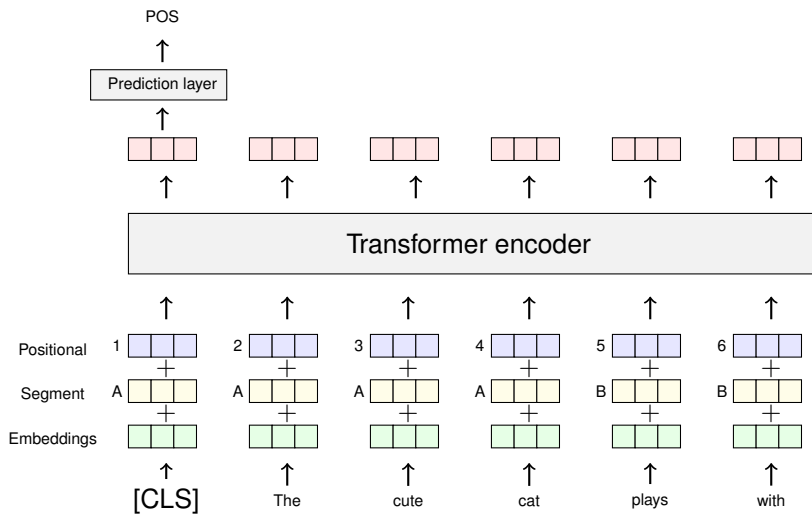
Through *finetuning* BERT, one can solve a number of language engineering problem.

- Token-based prediction (e.g. POS tagging) by adding a prediction layer on top of every output vector
- Sequence prediction (e.g. sentiment analysis) by adding a prediction layer on top of the CLS output vector

# Token classification



# Sequence classification



# Reading comprehension

Machine reading comprehension can be solved with the token classification approach.

**Input:** [CLS] *Wikipedia text about Stockholm* [SEP]  
Which scientific prize is awarded each year in Stockholm?

Categorize each token as

- B, first word of the answer, or
- I, word belongs to the answer, or
- O, word is not part of the answer



# BERT numbers

BERT Base/Bert Large):

- 12/24 stacked transformer blocks
- 12/16 attention heads
- 768/1024 hidden-vector size
- 512 input tokens (400 words)
- 235M/340M trainable parameters

Many larger/improved variants of BERT have been released 2019-

- More words in the input
- More training data
- Multilingual

# DD2417

## 9c: Autoregressive language models and GPT

Johan Boye, KTH

- GPT – Generative Pre-trained Transformer
- Developed by OpenAI
  - GPT 1–4 (2018–2022)
- Simpler training objective than BERT:
  - Predict the next token given the preceding ones
- GPT can also be fine-tuned to solve language tasks
  - Notably, InstructGPT, which formed the basis for ChatGPT

# Predicting the next token

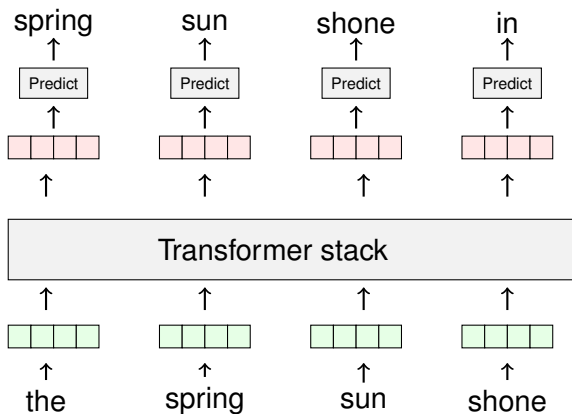
**Example:** The cute cat plays with a string

Training examples:


- cute **follows** The
- cat **follows** The cute
- plays **follows** The cute cat
- with **follows** The cute cat plays
- *etc.*

Using the Transformer architecture, all these examples can be learnt in parallel.

# Training an autoregressive language model



# Using an autoregressive language model



# Queries, keys, and values

$$y_1 = \alpha_{11} v_1 + \alpha_{12} v_2 + \alpha_{13} v_3 + \alpha_{14} v_4$$

$$y_2 = \alpha_{21} v_1 + \alpha_{22} v_2 + \alpha_{23} v_3 + \alpha_{24} v_4$$

$$y_3 = \alpha_{31} v_1 + \alpha_{32} v_2 + \alpha_{33} v_3 + \alpha_{34} v_4$$

$$y_4 = \alpha_{41} v_1 + \alpha_{42} v_2 + \alpha_{43} v_3 + \alpha_{44} v_4$$



	$k_1$	$k_2$	$k_3$	$k_4$
$q_1$	$\alpha_{11}$	$\alpha_{12}$	$\alpha_{13}$	$\alpha_{14}$
$q_2$	$\alpha_{21}$	$\alpha_{22}$	$\alpha_{23}$	$\alpha_{24}$
$q_3$	$\alpha_{31}$	$\alpha_{32}$	$\alpha_{33}$	$\alpha_{34}$
$q_4$	$\alpha_{41}$	$\alpha_{42}$	$\alpha_{43}$	$\alpha_{44}$

# Attention for autoregressive models

$$y_1 = \alpha_{11} v_1 + \alpha_{12} v_2 + \alpha_{13} v_3 + \alpha_{14} v_4$$

$$y_2 = \alpha_{21} v_1 + \alpha_{22} v_2 + \alpha_{23} v_3 + \alpha_{24} v_4$$

$$y_3 = \alpha_{31} v_1 + \alpha_{32} v_2 + \alpha_{33} v_3 + \alpha_{34} v_4$$

$$y_4 = \alpha_{41} v_1 + \alpha_{42} v_2 + \alpha_{43} v_3 + \alpha_{44} v_4$$

↑↑

	$k_1$	$k_2$	$k_3$	$k_4$
$q_1$	$\alpha_{11}$	$-\infty$	$-\infty$	$-\infty$
$q_2$	$\alpha_{21}$	$\alpha_{22}$	$-\infty$	$-\infty$
$q_3$	$\alpha_{31}$	$\alpha_{32}$	$\alpha_{33}$	$-\infty$
$q_4$	$\alpha_{41}$	$\alpha_{42}$	$\alpha_{43}$	$\alpha_{44}$



# GPT-2 numbers

- 12/24/36/48 stacked transformer blocks in GPT2
- 12/16/20/25 attention heads
- 768 hidden-vector size
- 1024 input tokens (800 words)
- 1.5B trainable parameters

It is assumed that...

- GPT-3 has 175B parameters
- GPT-4 has 1000B parameters
- GPT-3 and GPT-4 have a max input size of 4096 tokens (3200 words)